

How models fail — per-trial outcome composition by task and arm ·  $\geq 9B$ , think=off

